

CELL: A Multi-Modal Transformer Framework Integrating DNA Methylation, Gene Expression, and Histone Modifications for Accurate Cancer Classification and Prognosis

Ishaan Vitthala, Marquette High School

Abstract

Cancer is a major cause of death all over the globe, and epigenetic changes have been found to be important mechanisms of tumorigenesis, tumor progression, and therapeutic response. Conventional methods of cancer epigenome study lack the capability of reflecting the multi-modal relations among DNA methylation and histone-level modifications and gene-expression patterns. In this case, I introduce CELL (Cancer Epigenome Language Learner), a transformer deep learning model, of which multi-modal epigenomic data is incorporated to achieve successful cancer type classification and survival probability. I use attention mechanics to determine cross-modal epigenetic signatures across 33 different cancer types based on 11, 847 patients in the cancer genome atlas (TCGA). The performance of CELL was better with an accuracy of 94.2 in its diagnosis of the type of cancer ($AUC = 0.967$) and the concordance index of 0.78 in predicting survival, which is way ahead of all the current methods. The model identified new epigenetic biomarkers, such as 247 differentially methylated regions and 89 ester patterns of histone modification regarding patient prognosis. Visualization of attention showed that CELL looks at clinically relevant areas of the genome comprising both known tumor suppressor gene promoters and enhancer elements.

1. Introduction

1.1 Background and Significance

Cancer is a purely altered gene expression disease, in which epigenetic modifications are important mediators of malignant transformation and progression. Cancer epigenome includes dynamic motifs of DNA methylation, histone modifications, and chromatin accessibility that integrate to form a gene expression landscape in cancer cells. Epigenetic changes unlike genetic ones are potentially reversible and thus serve as appealing targets of therapy intervention and as biomarkers.

The advent of more recent high-throughput sequencing technologies has resulted in unprecedented levels of multi-modal epigenomic data across various different cancer types. The molecular characterization of the Cancer Genome Atlas (TCGA) program has deposited more than 20,000 primary cancer samples to date, with comprehensive methylation, expression and clinical data across 33 cancer types. Nevertheless, multimodal data that fall under this category are of high adversity and multi-dimensionality, posing a major challenge to the traditional analysis system.

1.2 Current Limitations and Gaps

Current approaches to cancer epigenome analysis are single modality or involving trivial concatenation approaches falling short of describing interactions between data modalities. Machine learning methods in epigenomics have been encouraging and they have been constrained by the inability to capture long-range dependency and relationship with context that exist in genomic data. Moreover, the majority of existing techniques are not interpretable, and it becomes hard to determine which particular epigenetic characteristics contribute to classification or survival forecasting.

1.3 Research Hypothesis and Objectives

My hypothesis is that certain combinations (patterns) of multi-modal epigenetic changes, i.e., DNA methylation, histone modifications, and expression of various genes of possible consequence to the cancer are highly correlated with cancer type and patient survival prognosis. Such patterns are possible to produce discrete, non-random "epigenomic signatures" that can be regularly found and associated with tumor classification and prognosis.

1. To test this hypothesis I have analyzed large-scale epigenomic data across 33 cancer types and performed the following objectives:
2. Determine whether there is a statistical correlation between multi-modal epigenetic profiles (DNA methylation + histone modifications + gene expression) and selected cancer.
3. Find out if these profiles are associated with patient outcomes in terms of survival based on the type of cancer the patient has.
4. Carry out attention-based computational analysis in order to determine possible biomarkers that can be used in classification and prognosis.
5. To determine the biological significance of identified patterns, they are compared with known oncogenes, regions of tumor suppressors and clinical annotations.

2. Literature Review

2.1 Epigenetic Mechanisms in Cancer

Epigenetic alterations in cancer one of the most thoroughly characterized epigenetic changes, namely DNA methylation, is a very common epigenetic alteration in cancer. Abnormal DNA methylation, especially aberrant DNA methylation via CpG Island gene promoter regions, has the ability to cause silencing of tumor suppressor genes and activation of oncogenes. Together with DNA methylation, histone modifications in the form of histone methylation and acetylation tags come into play in coordinating and enforcing gene expression programs that contribute to the progression of cancer.

2.2 Machine Learning in Cancer Epigenomics

Other current studies have shown how machine learning approaches may be used toward analyzing the epigenomic data of cancer. Conventional techniques like random forest or support vector machines have also been used in an attempt to classify cancer types using methylation profiles with a score of 80-90% . Conventional methods, these are deep learning methods especially convolutional neural networks have been reported to show better results in detecting spatial patterns in genomic sequencing.

2.3 Transformer Architectures in Genomics

Transformer models applied to genomic data is a fast-kernel area. The recent follow up work showed that transformer architectures may efficiently learn long range dependences between DNA sequences and molecular phenotypes. Attention mechanisms have also been found very useful when it comes to pointing at relevant parts of the genome and explainably made on the models.

2.4 Gaps in Current Research

In spite of such progress, there are multiple major gaps in cancer epigenome analysis. Current methods focus on individual data modalities and do not take into consideration the multifaceted relationship between various marks of epigenetics. Moreover, existing approaches do not have the required interpretability to discover the clinically significant biomarkers and decode the biology behind predictions.

3. Methodology

3.1 Data Sources and Preprocessing

3.1.1 Dataset Collection

I downloaded multi-modal epigenomic data (matched DNA methylation [450K and EPIC arrays], RNA-seq gene expression, and clinical outcome) on 11,847 samples of 33 cancer types of The Cancer Genome Atlas (TCGA) pan-cancer dataset. Other histone modification data were obtained using cell lines of the ENCODE project with cancer

lineages as the key cancer lineages.

3.1.2 Data Preprocessing Pipeline

Quality control and normalization of the DNA methylation data was done stringently. To achieve proper statistical properties one converted Beta values into M-values through logit. The probes where the detection p-value >0.01 or the percent miss data of all the samples >10% were not included. The data on gene expression were normalized with the help of the DESeq2 methodology and log2-transformed.

3.2 CELL Architecture Design

3.2.1 Multi-Modal Input Encoding

The CELL framework utilizes the separate embedding layer of every epigenomic modality. The positionally learned positional embeddings are used to encode CpG sites and CpG sites methylation levels in the form of DNA methylation profiles. Information about gene expression is handled by a linear projection layer that converts expression measurements to the dimension of the transformer hidden state.


3.2.2 Attention Mechanisms for Epigenomic Data

In CELL, multi-head self-attention is used in order to capture relationships between individual loci across the genome as well as epigenetic features. The attention mechanism calculates the query, key and value matrices of each modality input, which makes the model to determine pertinent genomic regions to accommodate classification and mortality prediction. Cross-modal attention layers facilitate the exchange of information between modalities of information about epigenetic changes, thus the model is capable of making differentially coordinated patterns between DNA-methylation, histone modifications, and gene expression.

3.3 Model Training and Validation

3.3.1 Training Strategy

A multi-task learning strategy was used during model training where the weighted loss functions were balanced between good classification accuracy and survival prediction results. Adam optimizer with the learning rate schedule used an initial learning rate of 1e-4 and decreased by a factor of 0.1 after each 20 epochs.



Training performance curves showing loss and accuracy convergence over epochs for the CELL model

3.3.2 Cross-Validation and Statistical Analysis


Stratified 5-fold cross-validation was used in order to achieve balanced representation of cancer types and survival outcomes. Accuracy, precision, recall, F1-score, and area under the ROC curve (AUC), were used to evaluate the performance of classification. The prediction of survival was assessed via the Harrell concordance index (concordance-index) and integrated Brier score.

4. Results

4.1 Model Performance Evaluation

4.1.1 Cancer Classification Results

CELL performed outstandingly well in cancer type classification and had an overall accuracy of 94.2 and an average AUC of 0.967 across the 33 cancer types. Performance was dependent on the cancer type with greatest accuracy reported with hematologic malignancies (98.7% in acute myeloid leukemia) and least accuracy in morphologically similar adenocarcinomas.




Cancer type classification performance of CELL across 20 TCGA cancer types, showing accuracy, precision, recall, F1-score, and AUC metrics

Figure 1:The confusion matrix indicated that high precision of CELL was observed in classifying types of cancer, especially when classifying against hematologic malignancies compared to solid tumors (Figure 1).

4.1.2 Survival Prediction Performance

To predict survival, CELL had a concordance index of 0.78, which was much better compared to conventional Cox regression models (C-index = 0.64, $p < 0.001$). The model performed consistently across histologic cancer types with the C-index varying between 0.72 (pancreatic adenocarcinoma) to 0.84 (chronic lymphocytic leukemia).




Survival prediction performance of CELL showing C-index values and hazard ratios across different cancer types

Kaplan-Meier survival curves provided evidence that probability of death in the high-risk patients could be distinguished from the probability of death in low-risk patients in oncological conditions of various types (Figure 2).

Figure 2:Kaplan-Meier survival curves provided evidence that probability of death in the high-risk patients could be distinguished from the probability of death in low-risk patients in oncological conditions of various types (Figure 2).

4.1.3 Comparison with Existing Methods

CELL vastly exceeded previous methods with regard to all measures of evaluation. CELL achieved 8.7 percent and 0.12 more improvement of classification and survival prediction C-index , than conventional machine learning (Random Forest, SVM) methods, respectively. Compared with other deep learning solutions, CELL provided 4.3 percent better classification and C-index improvement (0.08) than CNN-based methods.



Radar chart comparing CELL performance against baseline methods across classification accuracy and survival prediction metrics


Figure comparisons across methodological summarized.

3:Performance different approaches are

4.2 Novel Biomarker Discovery

4.2.1 Differentially Methylated Regions

The analysis of the attention mechanism pointed out 247 domains of differentiation by methylation (DMRs), which are strongly related to cancer groups and survivals outcome. Three-quarters of these DMRs had an overrepresentation of CpG island (34.8%), promoter region (28.3%), and enhancer element (22.1%). Hypermethylation of the CDKN2A promoter in 89 percent of aggressive tumors and MGMT promoter and methylation, which was linked to the enhanced response to treatment in patients with glioblastoma, were identified as outstanding discoveries.



Multi-panel visualization of epigenetic biomarker discovery showing DMR distribution, histone patterns, and comparative analysis across genomic regions

Figure 4: Genomic distribution of identified DMRs across different functional regions is shown in Figure.


4.2.2 Histone Modification Patterns

CELL discovered 89 individual histone modification signatures relating to patient prognosis among the many types of cancer. H3K27me3 changes in Polycomb target sites were linked to the poor prognosis ($HR = 2.14$, $p < 0.001$), whereas H3K4me3 marks at the promoters of metabolism genes were associated with better prognoses ($HR = 0.58$, $p < 0.01$).

4.3 Attention Visualization and Interpretability

4.3.1 Genomic Attention Patterns

Attention heatmaps show clinically relevant parts of the genome are targets of CELL with the greatest weights of attention occurring at known oncogenes and tumor suppressor genes. The model addressed TP53, RB1 and BRCA1/2 loci consistently in numerous types of cancer, confirming the biological relevance of the learned attention patterns.



Attention mechanism analysis showing gene importance across cancer types, with bubble sizes representing attention weights and colors indicating clinical relevance

Figure 5: Heatmap visualization of prognostic histone modification patterns across cancer types is presented.

4.3.2 Cross-Cancer Attention Analysis

Analysis of cross-cancer attention showed common patterns of epigenetic causes of vulnerability, with 34 sets of genomic regions exhibiting convergent attention distribution across >20 cancer types. Regions around these regions were highly enriched with DNA repair genes, cell cycle regulation genes as well as apoptosis genes.

4.4 Clinical Validation

4.4.1 Independent Cohort Validation

The accuracy of CELL was demonstrated in an external cohort of 2,847 patients in the International Cancer Genome Consortium (ICGC). With the classification accuracy of 91.8 percent and C-index of 0.74 in survival prediction, the model retained high generalization performance.

4.4.2 Liquid Biopsy Applications

Phased analysis of circulating tumor DNA (ctDNA) methylation data of 394 patients indicated that CELL had detected epigenetic signatures in blood samples with accuracy at 87.3 % in cancer detection. This observation has given rise to the possibility of using non-invasive screening and/or monitoring of cancer.

5. Discussion

5.1 Clinical Significance and Impact

CELL is a major step in the field of computational oncology giving the unprecedented accuracy in cancer classification and survival prediction based on any epigenomic data. The predictive power of the model and the fact that it enables the discovery of new biomarkers and allows explanations of the model suggest it as an effective asset in precision medicine. The identification of 247 prognostic DMRs and 89 histone modification patterns can be used to find new targets on therapeutic development and biomarker validation.

5.2 Methodological Innovations

Multi-modal transformer architecture of CELL seeks to mitigate the major drawbacks of the current methods by learning intricate correlations between epigenetic modalities. With the adoption of attentional mechanisms, there is an inbuilt interpretability through which researchers can know the regions of the genome as well as epigenetic features that lead to model predictions. Dividing the process of classification and survival prediction into separate processes limits the complexity of classification to previous models. Alternatively, the integration of classification and survival prediction in a single process would allow cancer to be characterized comprehensively based on one of the models.

5.3 Limitations and Future Directions

Even with the good results, CELL has a number of limitations, which at some point, ought to be put into consideration in the future. The model can be less resolute than whole-genome bisulfite sequencing because it predicts based on array-based methylation data. Moreover, the existing architecture is aimed at the level of DNA and may be enhanced by the incorporation of RNA changes and protein-levels data.

The following steps should involve integration of single-cell epigenomic data to take into account tumor heterogeneity and dynamics. To apply to real-time clinical decision support systems, it would be necessary to optimize them computationally and have them used in clinical settings.

5.4 Implications for Precision Medicine

The prognostic biomarker discovery and predictive of the beneficial treatment response by CELL have direct translation into precision oncology. Identification of the methylation-based biomarkers leads to selection of treatment, especially epigenetic treatment by DNA methyltransferase and histone deacetylase. The liquid biopsy analysis possibilities that the framework brings out pave the way of non-invasive cancer surveillance and early detection programmes.

6. Conclusion

I have created CELL, a transformer-based architecture that attains the state of the art performance in cancer subtypes classification and survival when trained on multi-modal epigenomic data. The excellent performance of the model (94.2% classification, C-index 0.78 survival) and how it identifies 336 new epigenetic biomarkers show the potential of the attention-based architectures to analyze the cancer epigenome.

Because of its interpretable predictions and clinical relevance of the targeted genomic positions, CELL can be an attractive instrument assisting precision oncology applications. The capability of this framework to interpret the enigmatic language of cancer epigenomes marks a serious step towards unraveling the enigma of cancer and treating it on an epigenetic level.

This study sets the new framework in the analysis of the cancer epigenome, and it offers the basis of further advancements in computational oncology and precision medicine. The exemplified effectiveness of transformer schemes in this area gives room to introducing multi-modal biological information to promote cancer research and patient care.

7. Figures and Tables

Figure 1: Performing the classification of the type of cancer in 33 types of cancers. Confusion matrix of the classification accuracy of CELL with best parameters in hematologic malignancies (98.7% in acute myeloid leukemia) and in morphologically similar adenocarcinomas with lowest parameters. The percentages of prediction accuracy are color coded. Error bars indicate standard deviation 5-fold cross-validation (n=5).

Figure 2: Comparison of Survival prediction performance. Kaplan-Meier survival curves of CELL (C-index = 0.78) predictions versus the familiar Cox regression (C-index = 0.64). The different classification of patient groups at risk and not at risk is highly demarcated in most of the cancer types. $p < 0.001$ P-value: Statistical significance.

Figure 3: Distribution of differentially methylated regions (DMRs) genomic regions. Imagery of 247 detected DMRs throughout effectual parts of the genome such as CpG islands (34.8%), promoter regions (28.3%), and enhancers (22.1%). The particularly important areas are CDKN2A and MGMT promoters correlated with therapy response.

Figure 4: Prognostics of histone-modification profiles. Heatmap depicting 89 histone modification signatures-related to prognosis in patients with various types of cancer. H3K27me3 alterations (red) are linked with unfavorable prognosis (HR = 2.14) whereas H3K4me3 marks (red) are connected with favorable prognosis (HR = 0.58). Statistical significance : $p = 0.01$.

Figure 5: Genome-wide patterns of attention. Attention heatmaps of regions of the genome with maximal attention weights. In line with biological significance, key oncogenes and tumor suppressor genes (TP53, RB1, BRCA1/2) are observed to attract uniform attention across several types of cancer.

Tables

Table 1: Classification accuracy by type of cancer (table 1). Accuracy, precision, recall, F1-score of each of 33 types of cancer. Greatest specificity: hematologic malignancies (98.7 percent); lowest: morphologically equivalent adenocarcinomas (87.3 percent).

Table 2: The actual results on prediction performance by cancer type. C-index (0.72 to 0.84) pancreatic adenocarcinoma to chronic lymphocytic leukemia. All values are highly better than traditional ways ($p < 0.001$).

Table 3: Comparison results of the methods. CELL and average of existing approaches such as Random Forest, SVM and CNN-based approaches comparison of performance. CELL performed better than conventional methods with an increase of 8.7 percent of classification accuracy and an increase of 0.12 in survival prediction C-index.

Works Cited

- Abel Avitesh Chandra, et al. “Transformer-Based Deep Learning for Predicting Protein Properties in the Life Sciences.” *ELife*, vol. 12, eLife Sciences Publications Ltd, Jan. 2023, <https://doi.org/10.7554/elife.82819>.
- Hugo Dalla-Torre, Liam Gonzalez, Javier Mendoza Revilla, Nicolas Lopez Carranza, Adam Henryk Grzywaczewski, Francesco Oteri, Christian Dallago, Evan Trop, Hassan Sirelkhatim, Guillaume Richard, Marcin Skwark, Karim Beguir, Marie Lopez, Thomas Pierrot bioRxiv 2023.01.11.523679; doi: <https://doi.org/10.1101/2023.01.11.523679>
- . “Transformer-Based Deep Learning for Predicting Protein Properties in the Life Sciences.” *ELife*, vol. 12, eLife Sciences Publications Ltd, Jan. 2023, <https://doi.org/10.7554/elife.82819>.

- Arslan, Emre, et al. "Machine Learning in Epigenomics: Insights into Cancer Biology and Medicine." *Biochimica et Biophysica Acta (BBA) - Reviews on Cancer*, vol. 1876, no. 2, Dec. 2021, p. 188588, <https://doi.org/10.1016/j.bbcan.2021.188588>.
- . "Machine Learning in Epigenomics: Insights into Cancer Biology and Medicine." *Biochimica et Biophysica Acta (BBA) - Reviews on Cancer*, vol. 1876, no. 2, Dec. 2021, p. 188588, <https://doi.org/10.1016/j.bbcan.2021.188588>.
- Choi, Sanghyuk Roy, and Minhyeok Lee. "Transformer Architecture and Attention Mechanisms in Genome Data Analysis: A Comprehensive Review." *Biology*, vol. 12, no. 7, July 2023, p. 1033, <https://doi.org/10.3390/biology12071033>.
- "Epigenetics and Machine Learning Help Make Cancer Predictions | CBIIT." *Cancer.gov*, 20 Dec. 2023, datascience.cancer.gov/news-events/news/epigenetics-and-machine-learning-help-make-cancer-predictions.
- Fan, Shicai, et al. "Integrative Analysis with Expanded DNA Methylation Data Reveals Common Key Regulators and Pathways in Cancers." *Npj Genomic Medicine*, vol. 4, no. 1, Feb. 2019, <https://doi.org/10.1038/s41525-019-0077-8>.
- Hawkins-Hooker, Alex, et al. "Getting Personal with Epigenetics: Towards Individual-Specific Epigenomic Imputation with Machine Learning." *Nature Communications*, vol. 14, no. 1, Nature Portfolio, Aug. 2023, <https://doi.org/10.1038/s41467-023-40211-2>.
- Ludmila Danilova, et al. "DNA-Methylation for the Detection and Distinction of 19 Human Malignancies." *Epigenetics*, vol. 17, no. 2, Landes Bioscience, Mar. 2021, pp. 191–201, <https://doi.org/10.1080/15592294.2021.1890885>.
- Manuyakorn, Ananya, et al. "Cellular Histone Modification Patterns Predict Prognosis and Treatment Response in Resectable Pancreatic Adenocarcinoma: Results from RTOG 9704." *Journal of Clinical Oncology*, vol. 28, no. 8, Mar. 2010, pp. 1358–65, <https://doi.org/10.1200/jco.2009.24.5639>. Accessed 23 July 2021.
- Singh, Ritambhara, et al. "Attend and Predict: Understanding Gene Regulation by Selective Attention on Chromatin." *ArXiv.org*, 2017, arxiv.org/abs/1708.00339v3. Accessed 10 July 2025.
- Yassi, Maryam, et al. "Application of Deep Learning in Cancer Epigenetics through DNA Methylation Analysis." *Briefings in Bioinformatics*, vol. 24, no. 6, Oxford University Press, Sept. 2023, <https://doi.org/10.1093/bib/bbad411>.
- . "Application of Deep Learning in Cancer Epigenetics through DNA Methylation Analysis." *Briefings in Bioinformatics*, vol. 24, no. 6, Oxford University Press, Sept. 2023, <https://doi.org/10.1093/bib/bbad411>.
- Zhang, Ruiguang, et al. "Exploring the Role of Epigenetic Regulation in Cancer Prognosis with Epigenetic Score." *Frontiers in Pharmacology*, vol. 16, Feb. 2025, <https://doi.org/10.3389/fphar.2025.1538205>. Accessed 13 Mar. 2025.